


Difficulty of sample-average value estimation in nonstationary problems


Charles C Onu, 260663256

Stationary environment setup



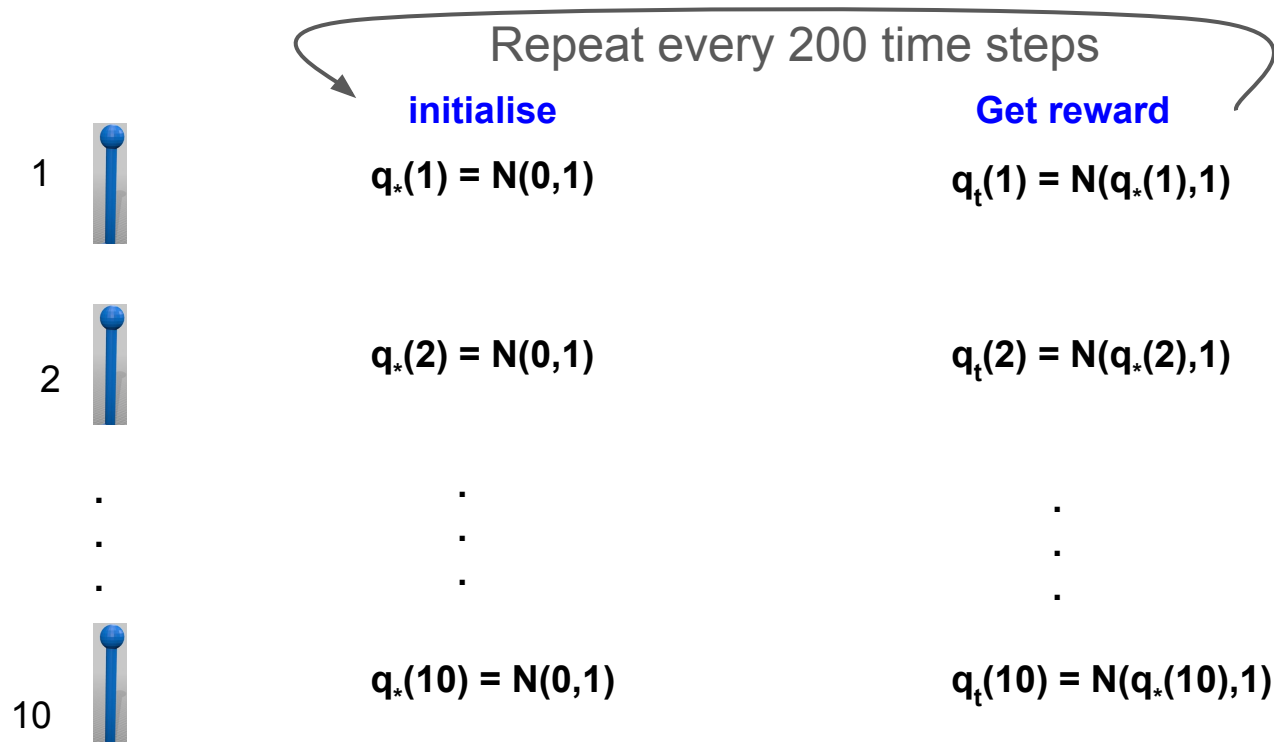
Agent runs for 1000 time steps ie $t = 1, \dots, 1000$

e-Greedy methods in stationary problem




NB: rewards are averaged across 2000 runs

Non stationary environment setup



Agent runs for 1000 time steps ie $t = 1, \dots, 1000$

Sample-averaging vs constant step-size for non stationary problem



Code

<https://github.com/rllabmcgill/rlcourse-january27-onucharles>